

L4-M8: Monitoring, Observability & Analytics - Logging Agent Actions, Observability Tools, KPIs for Agentic Systems

1. Introduction: The Imperative of Agentic Observability

The proliferation of autonomous AI agents, particularly those powered by Large Language Models (LLMs), has introduced a new layer of complexity to software systems. Unlike traditional monolithic or microservice architectures, agentic systems exhibit non-deterministic, multi-step reasoning, and tool-use behaviors, making their operational state notoriously opaque. **Observability**, a concept distinct from mere monitoring, becomes an essential control layer for these systems. It is the ability to infer the internal state of a system by examining its external outputs, a capability that is critical for debugging, performance optimization, and ensuring the safety and alignment of AI agents [1].

In the context of agentic systems, observability must extend beyond traditional infrastructure metrics to capture the unique aspects of the agent's decision-making process. This module delves into the technical foundations required to achieve this, focusing on the specialized logging of agent actions, the tools necessary to visualize these complex interactions, and the Key Performance Indicators (KPIs) that truly measure an agent's success and reliability.

2. The Pillars of Agentic Observability

Traditional system observability is built upon three foundational pillars: logs, metrics, and traces. For agentic systems, these pillars must be adapted and instrumented to

capture the *reasoning* and *intent* behind the agent's actions, not just the computational outcomes.

Pillar	Traditional Focus	Agentic Focus	Purpose in Agentic Systems
Logs	System events, errors, application output.	Agent state transitions, tool inputs/outputs, raw prompt/response pairs.	Debugging non-deterministic behavior, auditing, and compliance.
Metrics	CPU utilization, request latency, error rates.	Token usage, cost, task completion rate, tool call frequency, hallucination rate.	Quantifying performance, efficiency, and model quality.
Traces	Request flow across microservices.	The agent's "Chain of Thought" (CoT), decomposition of a task into sub-steps, reasoning path.	Understanding decision-making flow, identifying bottlenecks in the reasoning process.

2.1. The Challenge of Non-Determinism

One of the primary challenges in observing agentic systems is their inherent **non-determinism**. The same input may lead to different reasoning paths or final outputs due to variations in the LLM's stochastic nature (e.g., temperature setting), tool availability, or external environment changes. Robust observability provides the necessary data to analyze *why* a specific decision was made at a specific time, transforming an opaque black box into an auditable, transparent process [2].

3. Logging Agent Actions: Capturing the Agent's State

Effective logging for AI agents requires a structured approach that captures the entire context of an agent's operation. Simple print statements are insufficient; the logs must be machine-readable and contain rich metadata.

3.1. What to Log: Essential Agentic Data Points

A comprehensive agent log should capture the following critical data points at every significant step of the agent's execution loop:

- 1. Agent Identity and Context:** Agent ID, Session ID, User ID, Timestamp.
- 2. Input/Output:** Initial user prompt, final agent response.
- 3. LLM Interaction:**
 - **Prompt:** The exact, full prompt sent to the LLM (including system messages and history).
 - **Response:** The raw, unparsed response from the LLM.
 - **Metadata:** Model name, token usage (input/output), latency, cost.
- 4. Tool Use:**
 - **Tool Name:** The name of the tool selected by the agent.
 - **Tool Input:** The arguments passed to the tool.
 - **Tool Output:** The result returned by the tool.
- 5. State and Reasoning:**
 - **Internal State:** Key variables or memory content at the time of logging.
 - **Reasoning/CoT:** The intermediate "thought" or reasoning step generated by the LLM (if applicable).
 - **Action:** The specific action taken (e.g., `tool_call`, `final_answer`, `replan`).

3.2. Structured Logging and Standards

To facilitate analysis and integration with observability platforms, agent logs must adhere to a structured format, typically **JSON**. This allows log aggregation systems (like Elasticsearch, Splunk, or Loki) to index and query specific fields, such as `tool_name` or `token_usage`, efficiently.

OpenTelemetry (OTel), the industry standard for telemetry data, provides a unified specification for logs, metrics, and traces. While OTel is primarily known for tracing, its logging specification encourages the inclusion of trace and span IDs within the log record, effectively linking a specific log message to its place in the overall execution trace [3].

```
{  
    "timestamp": "2025-10-30T10:00:00Z",  
    "level": "INFO",  
    "service.name": "financial-agent-v2",  
    "trace_id": "a1b2c3d4e5f67890",  
    "span_id": "1234567890abcdef",  
    "agent.action": "tool_call",  
    "user.id": "user-456",  
    "tool.name": "stock_price_api",  
    "tool.input": {  
        "ticker": "GOOGL",  
        "date": "2025-10-29"  
    },  
    "llm.metadata": {  
        "model": "gpt-4.1-turbo",  
        "prompt_tokens": 120,  
        "completion_tokens": 5  
    },  
    "message": "Agent called stock_price_api for GOOGL."  
}
```

This structured log is far more valuable than a simple text string, as it allows for complex filtering and aggregation, such as calculating the average token usage for all tool calls to `stock_price_api`.

3.3. Step-by-Step: Implementing Structured Logging in Python

A practical implementation involves using a standard logging library (like Python's `logging`) in conjunction with a structured logging formatter (like `python-json-logger`).

Step 1: Setup

```
pip install python-json-logger
```

Step 2: Configuration and Logging

```

import logging
import json_log_formatter

# Configure the logger to use the JSON formatter
formatter = json_log_formatter.JSONFormatter()
handler = logging.StreamHandler()
handler.setFormatter(formatter)

logger = logging.getLogger('agent_logger')
logger.addHandler(handler)
logger.setLevel(logging.INFO)

def log_agent_tool_call(tool_name, tool_input, llm_metadata, trace_id,
span_id):
    """Logs a structured event for a tool call."""
    extra = {
        'trace_id': trace_id,
        'span_id': span_id,
        'agent.action': 'tool_call',
        'tool.name': tool_name,
        'tool.input': tool_input,
        'llm.metadata': llm_metadata
    }
    logger.info("Tool call initiated.", extra=extra)

# Example Usage within an agent loop
current_trace_id = "a1b2c3d4e5f67890"
current_span_id = "1234567890abcdef"
metadata = {"model": "gemini-2.5-flash", "prompt_tokens": 150}

log_agent_tool_call(
    tool_name="weather_api",
    tool_input={"city": "London"},
    llm_metadata=metadata,
    trace_id=current_trace_id,
    span_id=current_span_id
)

```

4. Tracing the Agent's Chain of Thought

While logs capture individual events, **tracing** is essential for visualizing the temporal and causal relationships between these events. In agentic systems, a trace represents the entire execution of a user query, and the individual spans within that trace map directly to the agent's internal reasoning steps.

4.1. The Need for Tracing in LLM Pipelines

A typical agent execution involves a sequence of steps: 1. **Initial Prompt Processing** 2. **Reasoning Step 1** (e.g., deciding to call a tool) 3. **Tool Execution** 4. **Reasoning Step 2** (e.g., synthesizing tool output into a final answer) 5. **Final Response Generation**

A traditional log file would show these events chronologically, but it would not easily illustrate the *hierarchy* or the *latency* of each step. Tracing, using a standard like OpenTelemetry, provides this crucial context by organizing events into nested **spans**.

Span Attribute	Description	Example Value
llm.request.type	The type of LLM call.	chat, completion, embedding
agent.step.type	The nature of the agent's action.	reasoning, tool_call, replan
tool.name	The specific tool invoked.	database_query_tool
db.statement	The query executed by a database tool.	SELECT * FROM users WHERE id=1
agent.prompt.hash	A hash of the full prompt for de-duplication.	0xABCD123DEF456

4.2. OpenTelemetry for Agent Tracing

OpenTelemetry (OTel) provides language-agnostic APIs and SDKs to instrument code. For an AI agent, the instrumentation involves wrapping the core components of the agent framework (e.g., LangChain, LlamaIndex, or custom loops) to automatically create, start, and end spans.

Step-by-step: Instrumenting an Agent with OTel

- 1. Setup OTel SDK and Exporter:** Install the necessary OTel packages (e.g., `opentelemetry-api`, `opentelemetry-sdk`, and an exporter like `opentelemetry-exporter-otlp`).
- 2. Initialize Tracer Provider:** Configure the OTel SDK to send traces to a collector (e.g., Jaeger, Tempo).
- 3. Define the Agent's Process as a Root Span:** The entire execution of the agent for a single user query is the root span.

```
```python
from opentelemetry import trace
tracer = trace.get_tracer("agent.system")```

```

```
def run_agent(user_query):
 with tracer.start_as_current_span("agent-execution",
 attributes={"user.query": user_query}) as root_span:
 # ... start agent```

```

```
logic ... process_query(user_query) 4. **Create Nested Spans for Internal
Steps:** Every significant step, such as an LLM call or a tool
invocation, is wrapped in its own span, nested under the root
span. python def decide_action(context): with
tracer.start_as_current_span("llm-reasoning") as span: # Call LLM to decide next
step llm_response = call_llm(context) span.set_attribute("llm.model", "gpt-4.1-
turbo") span.set_attribute("llm.prompt_tokens", len(context.tokens)) return
parse_action(llm_response) ``
```

## 5. Inject Context for

**Asynchronous/Distributed Agents:** For multi-agent systems or agents that interact with separate services, **context propagation** is vital. OTel automatically handles injecting the current trace context (Trace ID and Span ID) into network requests, ensuring that all logs and traces generated by downstream services are correctly linked back to the original root trace.

## 5. Observability Tools and Platforms

---

The data generated by instrumented agents—structured logs, metrics, and traces—must be collected, stored, and visualized. A modern observability stack is essential for transforming raw telemetry into actionable insights.

### 5.1. The Modern Observability Stack (LGT)

A common, open-source stack for observability is the **Prometheus, Grafana, and Loki/Tempo** combination, often referred to as the LGT stack.

Tool	Function	Application in Agentic Systems
Prometheus	Time-series database and metric collection.	Scrapes agent metrics (e.g., P95 latency, tool success rate) via an exposed <code>/metrics</code> endpoint.
Grafana	Data visualization and dashboarding.	Creates dashboards to monitor agent health, performance KPIs, and cost trends over time.
Loki	Log aggregation system (optimized for logs).	Indexes and stores structured JSON logs, allowing powerful queries based on log labels and content.
Tempo	Distributed tracing backend.	Stores and queries OpenTelemetry traces, enabling visualization of the agent's CoT flow.

## 5.2. Practical Application: A Grafana Dashboard for Agent Health

A well-designed Grafana dashboard for an agentic system would typically include:

- **Service Level Objectives (SLOs):** A gauge showing the percentage of tasks completed within the target latency (e.g., 99% of tasks complete in < 5 seconds).
- **Cost Monitoring:** A graph tracking daily token usage and estimated cost per agent, broken down by model.
- **Tool Reliability:** A table showing the error rate and average latency for each external tool the agent uses.
- **Reasoning Depth:** A histogram of the number of steps (spans) per trace, indicating the complexity of the tasks being handled.
- **Log Explorer:** An embedded Loki panel allowing operators to filter logs by `agent.action` or `tool.name` to quickly drill down into failures identified on the metrics panels.

## 6. Key Performance Indicators (KPIs) for Agentic Systems

Traditional KPIs like server uptime and request throughput are insufficient for AI agents. Agentic systems require a set of metrics that quantify the quality of the agent's

*output* and the *efficiency* of its reasoning process. These KPIs fall into four main categories: Performance, Quality, Cost, and Safety/Compliance [4].

## 6.1. Performance and Efficiency KPIs

These metrics measure how quickly and efficiently the agent achieves its goals.

- **Task Completion Rate (TCR):** The percentage of user requests that result in a successful, final, and correct answer.
  - *Formula:*  $(\text{Successful Tasks} / \text{Total Tasks}) * 100$
- **Average Steps to Completion (ASC):** The mean number of reasoning steps (spans) required to complete a task. A high ASC may indicate inefficient planning or excessive replanning.
- **P95/P99 Latency:** The time taken for 95% or 99% of tasks to complete. This is a critical Service Level Indicator (SLI).
- **Tool Call Success Rate:** The percentage of tool calls that return a valid, non-error response.

## 6.2. Quality and Alignment KPIs

These metrics are crucial for ensuring the agent's output is valuable and aligned with user intent.

- **Hallucination Rate:** The frequency with which the agent generates factually incorrect or unsupported information. This is often measured via a separate, automated validation step or human feedback.
- **Rejection Rate / User Feedback Score (UFS):** The percentage of agent responses explicitly rejected or rated poorly by the user.
- **Tool Selection Accuracy:** The percentage of times the agent correctly identifies and calls the most appropriate tool for a given sub-task.
- **Response Relevance Score:** A score, often generated by a separate LLM or human evaluator, measuring how well the final answer addresses the original user prompt.

### 6.3. Cost and Resource KPIs

Monitoring resource consumption is vital for financial control, especially with pay-per-token models.

- **Average Cost Per Task (ACPT):** The total cost (LLM API calls + compute) divided by the number of completed tasks.
- **Token Efficiency (TE):** The ratio of output tokens to input tokens, or the ratio of total tokens to task complexity. High token usage for simple tasks indicates inefficiency.
- **Cache Hit Rate:** For systems utilizing caching (e.g., for common tool calls or prompts), this measures the percentage of requests served from the cache, reducing both latency and cost.

### 6.4. Safety and Compliance KPIs

These are non-functional requirements that ensure the agent operates within defined boundaries.

- **Guardrail Violation Rate:** The frequency of attempts to bypass safety filters or generate harmful/unaligned content.
- **PII Exposure Count:** The number of log entries or responses containing detected Personally Identifiable Information (PII).
- **Access Control Failure Rate:** The frequency of unauthorized tool or data access attempts.

## 7. Real-World Example: Debugging a Financial Agent Failure

---

Consider a **Financial Analyst Agent** designed to answer complex investment questions using a portfolio management tool and a real-time stock data API.

**Scenario:** A user asks, "What was the return on my GOOGL stock in Q3 2025?" The agent returns an error: "Could not calculate return. Data source failed."

Log/Trace Data Point	Value	Insight
Root Trace ID	trace-xyz-123	Links all subsequent events.
Span 1 ( <code>llm-reasoning</code> )	Duration: 150ms. Output: <pre>ToolCall(name='portfolio_tool', args={'ticker': 'GOOGL', 'period': 'Q3 2025'})</pre>	Agent correctly identified the need for the portfolio tool. Reasoning was fast.
Span 2 ( <code>portfolio_tool</code> )	Duration: 4,500ms. Status: <code>ERROR</code> . Log: <pre>Tool call failed: 'stock_api_v2' endpoint timed out after 4s.</pre>	<b>Key Insight:</b> The tool call itself failed, but the <i>reason</i> was a dependency (the stock API) timing out. The agent's error message was too generic.
Metric: <code>tool_call_latency_p99</code>	5,100ms (for <code>portfolio_tool</code> )	<b>Systemic Issue:</b> The P99 latency for this tool is consistently high, suggesting a chronic performance issue with the underlying service.
Log: <code>agent.action: replan</code>	Count: 0	<b>Agent Limitation:</b> The agent did not attempt to replan or use an alternative tool (e.g., a cached data source) after the failure, indicating a lack of robustness.

**Step-by-Step Debugging:** 1. **Monitor:** The Grafana dashboard shows a spike in the `Tool Call Failure Rate` metric for the `portfolio_tool`. 2. **Trace:** An operator selects a failed trace (`trace-xyz-123`) from the Tempo UI. The trace visualization immediately highlights the `portfolio_tool` span as the bottleneck and failure point.

3. **Log Drill-down:** By clicking on the failed span, the operator is directed to the associated structured log in Loki, which reveals the specific timeout error from the *internal* dependency (`stock_api_v2`). 4. **Action:** The team identifies the `stock_api_v2` service as the root cause, not the agent's logic. They also update the agent's prompt to include a **re-planning instruction** for tool failures, improving its resilience.

This example demonstrates how the combination of logs (specific error message), traces (causal flow and latency), and metrics (systemic failure rate) provides a complete picture of the failure, something impossible with just one data source.

## 8. Conclusion and Key Takeaways

---

Observability is not an optional feature but a fundamental requirement for deploying reliable, cost-effective, and safe AI agentic systems. The non-deterministic nature of these systems necessitates a shift from traditional monitoring to a comprehensive strategy that captures the agent's internal reasoning and state transitions.

The technical foundation for agentic observability rests on: 1. **Structured Logging:** Capturing rich, machine-readable data points (prompts, tool inputs/outputs, state) at every step of the agent's execution. 2. **Distributed Tracing:** Using standards like **OpenTelemetry** to visualize the agent's complex Chain of Thought (CoT) as a series of nested spans, allowing for precise latency analysis and bottleneck identification. 3. **Specialized KPIs:** Moving beyond infrastructure metrics to focus on **Task Completion Rate**, **Token Efficiency**, and **Hallucination Rate**, which directly measure the agent's business value and alignment.

By implementing a robust observability stack (e.g., Prometheus, Grafana, Loki, Tempo) and instrumenting agents with OTel, organizations can gain the necessary transparency to debug, optimize, and ultimately trust their autonomous AI workforce.

---

## References

---

- [1] IBM. *Why observability is essential for AI agents*. Available at: <https://www.ibm.com/think/insights/ai-agent-observability>
- [2] Microsoft Azure. *Top 5 agent observability best practices for reliable AI*. Available at:

<https://azure.microsoft.com/en-us/blog/agent-factory-top-5-agent-observability-best-practices-for-reliable-ai/> [3] Uptrace. *OpenTelemetry for AI Systems: Implementation Guide*. Available at: <https://uptrace.dev/blog/opentelemetry-ai-systems> [4] Auxiliobits. *Evaluating Agentic AI in the Enterprise: Metrics, KPIs, and Benchmarks*. Available at: <https://www.auxiliobits.com/blog/evaluating-agentic-ai-in-the-enterprise-metrics-kpis-and-benchmarks/> [5] Galileo AI. *Master Logging and Tracing for Effective AI Development*. Available at: <https://galileo.ai/blog/logging-tracing-ai-systems>